

In [1]:

```
# Importing the Libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
```

Loading the Dataset

In [2]:

```
from sklearn.datasets import load_boston
boston =load_boston()
boston.keys()
```

Out[2]:

```
dict_keys(['data', 'target', 'feature_names', 'DESCR', 'filename'])
```

Dataset Description

In [3]:

```
print(boston.DESCR)
```

.. _boston_dataset:

Boston house prices dataset

Data Set Characteristics:

:Number of Instances: 506

:Number of Attributes: 13 numeric/categorical predictive. Median Value (attribute 14) is usually the target.

:Attribute Information (in order):

- CRIM per capita crime rate by town
- ZN proportion of residential land zoned for lots over 25,000 sq.ft.

- INDUS proportion of non-retail business acres per town
- CHAS Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)

- NOX nitric oxides concentration (parts per 10 million)
- RM average number of rooms per dwelling

- AGE proportion of owner-occupied units built prior to 1940

- DIS weighted distances to five Boston employment centres

- RAD index of accessibility to radial highways

- TAX full-value property-tax rate per \$10,000

- PTRATIO pupil-teacher ratio by town

- B $1000(Bk - 0.63)^2$ where Bk is the proportion of blacks by town

- LSTAT % lower status of the population

- MEDV Median value of owner-occupied homes in \$1000's

:Missing Attribute Values: None

:Creator: Harrison, D. and Rubinfeld, D.L.

This is a copy of UCI ML housing dataset.

<https://archive.ics.uci.edu/ml/machine-learning-databases/housing/>

This dataset was taken from the StatLib library which is maintained at Carnegie Mellon University.

The Boston house-price data of Harrison, D. and Rubinfeld, D.L. 'Hedonic prices and the demand for clean air', J. Environ. Economics & Management, vol.5, 81-102, 1978. Used in Belsley, Kuh & Welsch, 'Regression diagnostics

...', Wiley, 1980. N.B. Various transformations are used in the table on pages 244-261 of the latter.

The Boston house-price data has been used in many machine learning papers that address regression problems.

.. topic:: References

- Belsley, Kuh & Welsch, 'Regression diagnostics: Identifying Influential Data and Sources of Collinearity', Wiley, 1980. 244-261.

- Quinlan,R. (1993). Combining Instance-Based and Model-Based Learning. In Proceedings on the Tenth International Conference of Machine Learning, 236-243, University of Massachusetts, Amherst. Morgan Kaufmann.

In [4]:

```
print(boston.data)
```

```
[[6.3200e-03 1.8000e+01 2.3100e+00 ... 1.5300e+01 3.9690e+02 4.9800e+00]
 [2.7310e-02 0.0000e+00 7.0700e+00 ... 1.7800e+01 3.9690e+02 9.1400e+00]
 [2.7290e-02 0.0000e+00 7.0700e+00 ... 1.7800e+01 3.9283e+02 4.0300e+00]
 ...
 [6.0760e-02 0.0000e+00 1.1930e+01 ... 2.1000e+01 3.9690e+02 5.6400e+00]
 [1.0959e-01 0.0000e+00 1.1930e+01 ... 2.1000e+01 3.9345e+02 6.4800e+00]
 [4.7410e-02 0.0000e+00 1.1930e+01 ... 2.1000e+01 3.9690e+02 7.8800e+00]]
```

In [5]:

```
print(boston.target)
```

```
[24. 21.6 34.7 33.4 36.2 28.7 22.9 27.1 16.5 18.9 15. 18.9 21.7 20.4
 18.2 19.9 23.1 17.5 20.2 18.2 13.6 19.6 15.2 14.5 15.6 13.9 16.6 14.8
 18.4 21. 12.7 14.5 13.2 13.1 13.5 18.9 20. 21. 24.7 30.8 34.9 26.6
 25.3 24.7 21.2 19.3 20. 16.6 14.4 19.4 19.7 20.5 25. 23.4 18.9 35.4
 24.7 31.6 23.3 19.6 18.7 16. 22.2 25. 33. 23.5 19.4 22. 17.4 20.9
 24.2 21.7 22.8 23.4 24.1 21.4 20. 20.8 21.2 20.3 28. 23.9 24.8 22.9
 23.9 26.6 22.5 22.2 23.6 28.7 22.6 22. 22.9 25. 20.6 28.4 21.4 38.7
 43.8 33.2 27.5 26.5 18.6 19.3 20.1 19.5 19.5 20.4 19.8 19.4 21.7 22.8
 18.8 18.7 18.5 18.3 21.2 19.2 20.4 19.3 22. 20.3 20.5 17.3 18.8 21.4
 15.7 16.2 18. 14.3 19.2 19.6 23. 18.4 15.6 18.1 17.4 17.1 13.3 17.8
 14. 14.4 13.4 15.6 11.8 13.8 15.6 14.6 17.8 15.4 21.5 19.6 15.3 19.4
 17. 15.6 13.1 41.3 24.3 23.3 27. 50. 50. 50. 22.7 25. 50. 23.8
 23.8 22.3 17.4 19.1 23.1 23.6 22.6 29.4 23.2 24.6 29.9 37.2 39.8 36.2
 37.9 32.5 26.4 29.6 50. 32. 29.8 34.9 37. 30.5 36.4 31.1 29.1 50.
 33.3 30.3 34.6 34.9 32.9 24.1 42.3 48.5 50. 22.6 24.4 22.5 24.4 20.
 21.7 19.3 22.4 28.1 23.7 25. 23.3 28.7 21.5 23. 26.7 21.7 27.5 30.1
 44.8 50. 37.6 31.6 46.7 31.5 24.3 31.7 41.7 48.3 29. 24. 25.1 31.5
 23.7 23.3 22. 20.1 22.2 23.7 17.6 18.5 24.3 20.5 24.5 26.2 24.4 24.8
 29.6 42.8 21.9 20.9 44. 50. 36. 30.1 33.8 43.1 48.8 31. 36.5 22.8
 30.7 50. 43.5 20.7 21.1 25.2 24.4 35.2 32.4 32. 33.2 33.1 29.1 35.1
 45.4 35.4 46. 50. 32.2 22. 20.1 23.2 22.3 24.8 28.5 37.3 27.9 23.9
 21.7 28.6 27.1 20.3 22.5 29. 24.8 22. 26.4 33.1 36.1 28.4 33.4 28.2
 22.8 20.3 16.1 22.1 19.4 21.6 23.8 16.2 17.8 19.8 23.1 21. 23.8 23.1
 20.4 18.5 25. 24.6 23. 22.2 19.3 22.6 19.8 17.1 19.4 22.2 20.7 21.1
 19.5 18.5 20.6 19. 18.7 32.7 16.5 23.9 31.2 17.5 17.2 23.1 24.5 26.6
 22.9 24.1 18.6 30.1 18.2 20.6 17.8 21.7 22.7 22.6 25. 19.9 20.8 16.8
 21.9 27.5 21.9 23.1 50. 50. 50. 50. 13.8 13.8 15. 13.9 13.3
 13.1 10.2 10.4 10.9 11.3 12.3 8.8 7.2 10.5 7.4 10.2 11.5 15.1 23.2
 9.7 13.8 12.7 13.1 12.5 8.5 5. 6.3 5.6 7.2 12.1 8.3 8.5 5.
 11.9 27.9 17.2 27.5 15. 17.2 17.9 16.3 7. 7.2 7.5 10.4 8.8 8.4
 16.7 14.2 20.8 13.4 11.7 8.3 10.2 10.9 11. 9.5 14.5 14.1 16.1 14.3
 11.7 13.4 9.6 8.7 8.4 12.8 10.5 17.1 18.4 15.4 10.8 11.8 14.9 12.6
 14.1 13. 13.4 15.2 16.1 17.8 14.9 14.1 12.7 13.5 14.9 20. 16.4 17.7
 19.5 20.2 21.4 19.9 19. 19.1 19.1 20.1 19.9 19.6 23.2 29.8 13.8 13.3
 16.7 12. 14.6 21.4 23. 23.7 25. 21.8 20.6 21.2 19.1 20.6 15.2 7.
 8.1 13.6 20.1 21.8 24.5 23.1 19.7 18.3 21.2 17.5 16.8 22.4 20.6 23.9
 22. 11.9]
```

In [8]:

```
print(boston.feature_names)
```

```
['CRIM' 'ZN' 'INDUS' 'CHAS' 'NOX' 'RM' 'AGE' 'DIS' 'RAD' 'TAX' 'PTRATIO'
 'B' 'LSTAT']
```

In [9]:

```
# Prepare the dataframe
df = pd.DataFrame(boston.data, columns=boston.feature_names)
```

In [10]:

```
df.head()
```

Out[10]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90

In [11]:

```
df['Price']=boston.target
```

In [12]:

```
df.head()
```

Out[12]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90

In [18]:

```
# To check the shape of the dataset
df.shape
```

Out[18]:

(506, 14)

In [13]:

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
 #   Column   Non-Null Count  Dtype  
--- 
 0   CRIM      506 non-null    float64
 1   ZN         506 non-null    float64
 2   INDUS     506 non-null    float64
 3   CHAS       506 non-null    float64
 4   NOX        506 non-null    float64
 5   RM          506 non-null    float64
 6   AGE         506 non-null    float64
 7   DIS          506 non-null    float64
 8   RAD          506 non-null    float64
 9   TAX          506 non-null    float64
 10  PTRATIO    506 non-null    float64
 11  B           506 non-null    float64
 12  LSTAT      506 non-null    float64
 13  Price       506 non-null    float64
dtypes: float64(14)
memory usage: 55.5 KB
```

In [14]:

```
# Statistical information for the data
df.describe().T
```

Out[14]:

	count	mean	std	min	25%	50%	75%
CRIM	506.0	3.613524	8.601545	0.00632	0.082045	0.25651	3.677083
ZN	506.0	11.363636	23.322453	0.00000	0.000000	0.00000	12.500000
INDUS	506.0	11.136779	6.860353	0.46000	5.190000	9.69000	18.100000
CHAS	506.0	0.069170	0.253994	0.00000	0.000000	0.00000	0.000000
NOX	506.0	0.554695	0.115878	0.38500	0.449000	0.53800	0.624000
RM	506.0	6.284634	0.702617	3.56100	5.885500	6.20850	6.623500
AGE	506.0	68.574901	28.148861	2.90000	45.025000	77.50000	94.075000
DIS	506.0	3.795043	2.105710	1.12960	2.100175	3.20745	5.188425
RAD	506.0	9.549407	8.707259	1.00000	4.000000	5.00000	24.000000
TAX	506.0	408.237154	168.537116	187.00000	279.000000	330.00000	666.000000
PTRATIO	506.0	18.455534	2.164946	12.60000	17.400000	19.05000	20.200000
B	506.0	356.674032	91.294864	0.32000	375.377500	391.44000	396.225000
LSTAT	506.0	12.653063	7.141062	1.73000	6.950000	11.36000	16.955000
Price	506.0	22.532806	9.197104	5.00000	17.025000	21.20000	25.000000

In [15]:

```
## Checking the Missing value
df.isnull().sum()
```

Out[15]:

```
CRIM      0
ZN        0
INDUS     0
CHAS      0
NOX       0
RM        0
AGE       0
DIS       0
RAD       0
TAX       0
PTRATIO   0
B          0
LSTAT     0
Price     0
dtype: int64
```

Exploratory Data Analysis

In [16]:

```
# Checking Correlation between its Features
df.corr()
```

Out[16]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	D
CRIM	1.000000	-0.200469	0.406583	-0.055892	0.420972	-0.219247	0.352734	-0.3796
ZN	-0.200469	1.000000	-0.533828	-0.042697	-0.516604	0.311991	-0.569537	0.6644
INDUS	0.406583	-0.533828	1.000000	0.062938	0.763651	-0.391676	0.644779	-0.7080
CHAS	-0.055892	-0.042697	0.062938	1.000000	0.091203	0.091251	0.086518	-0.0991
NOX	0.420972	-0.516604	0.763651	0.091203	1.000000	-0.302188	0.731470	-0.7692
RM	-0.219247	0.311991	-0.391676	0.091251	-0.302188	1.000000	-0.240265	0.2052
AGE	0.352734	-0.569537	0.644779	0.086518	0.731470	-0.240265	1.000000	-0.7478
DIS	-0.379670	0.664408	-0.708027	-0.099176	-0.769230	0.205246	-0.747881	1.0000
RAD	0.625505	-0.311948	0.595129	-0.007368	0.611441	-0.209847	0.456022	-0.4945
TAX	0.582764	-0.314563	0.720760	-0.035587	0.668023	-0.292048	0.506456	-0.5344
PTRATIO	0.289946	-0.391679	0.383248	-0.121515	0.188933	-0.355501	0.261515	-0.2324
B	-0.385064	0.175520	-0.356977	0.048788	-0.380051	0.128069	-0.273534	0.2915
LSTAT	0.455621	-0.412995	0.603800	-0.053929	0.590879	-0.613808	0.602339	-0.4969
Price	-0.388305	0.360445	-0.483725	0.175260	-0.427321	0.695360	-0.376955	0.2499

Data Visualization- Performing basic Exploratory operation related to feature

In [17]:

```
import seaborn as sns
sns.set(rc={'figure.figsize':(15,10)})
sns.heatmap(df.corr(), annot=True)
```

Out[17]:

<AxesSubplot: >

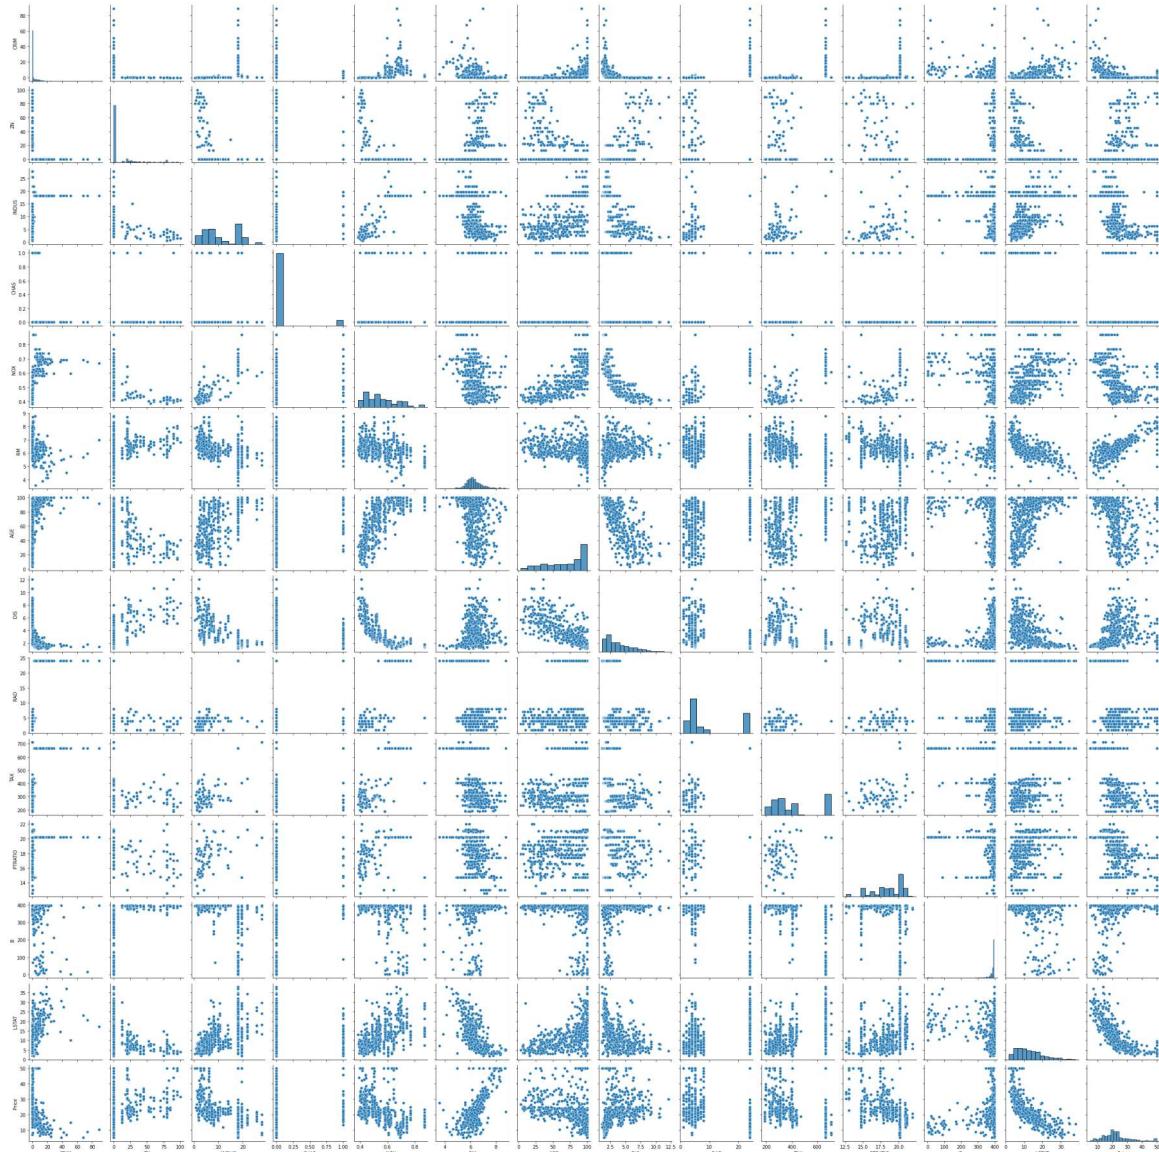


In [19]:

```
# Pairplot
sns.pairplot(df)
```

Out[19]:

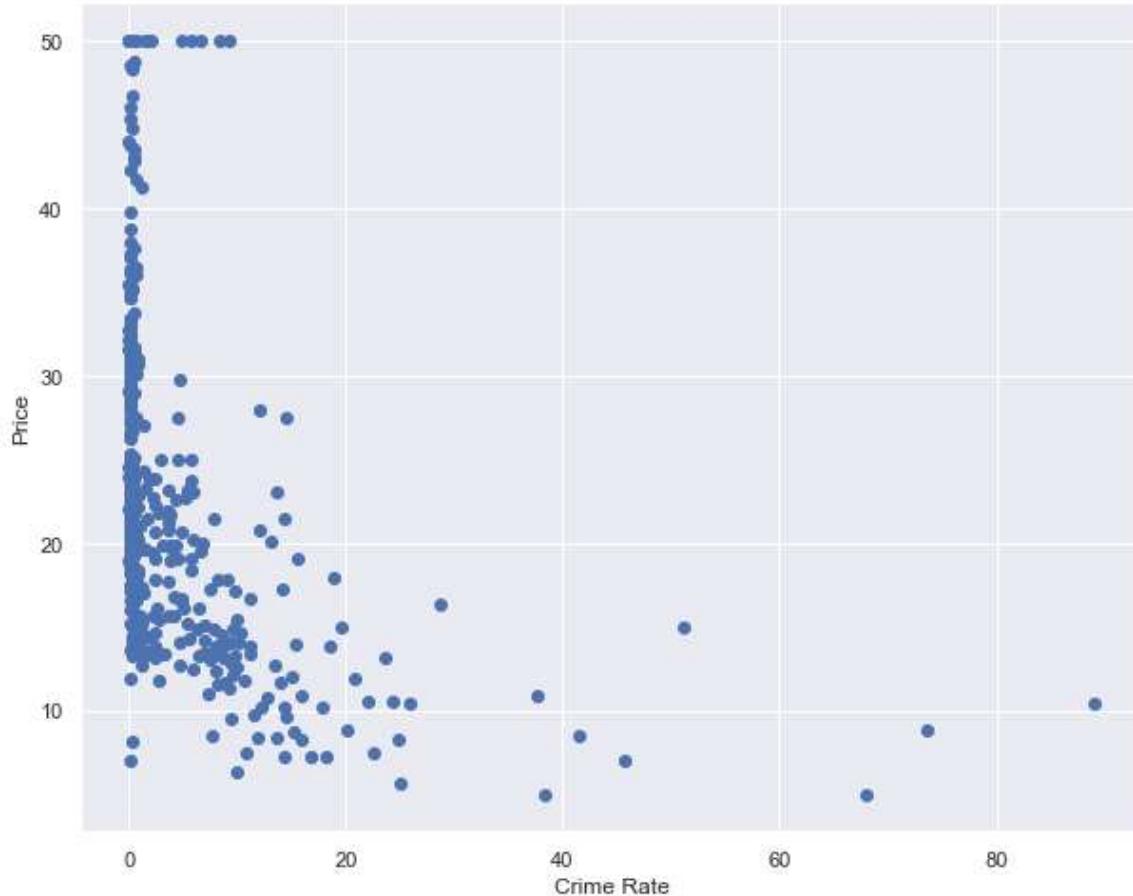
```
<seaborn.axisgrid.PairGrid at 0x1e73403bbb0>
```



Scatterplot - Crime Rate Vs Price

In [20]:

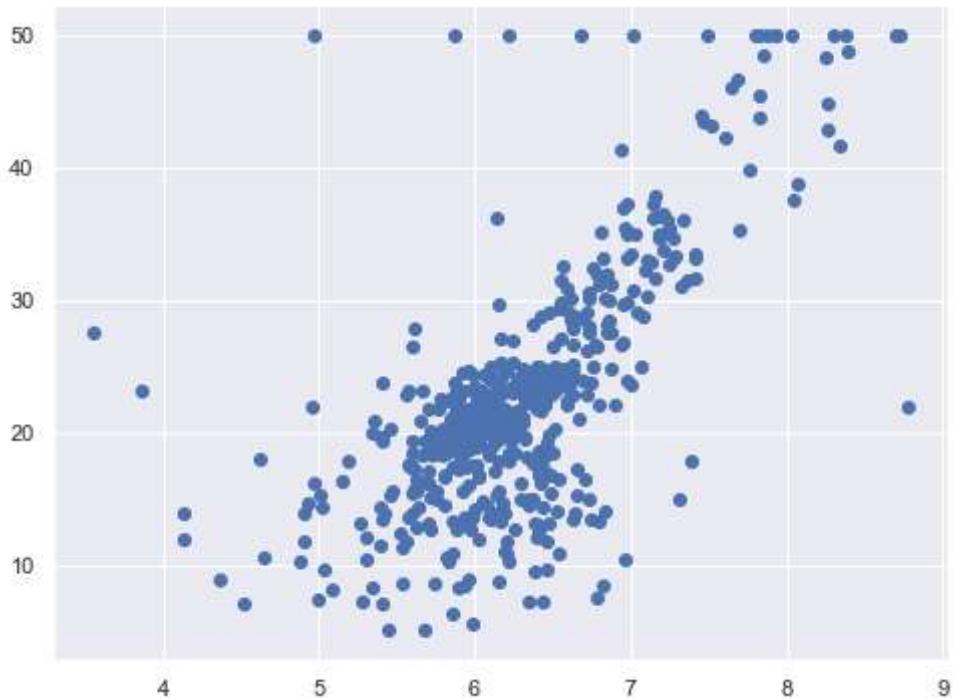
```
sns.set(rc={'figure.figsize':(10,8)})  
plt.scatter(df['CRIM'],df['Price'])  
plt.xlabel('Crime Rate')  
plt.ylabel('Price');
```



Scatterplot - Rooms Vs Price

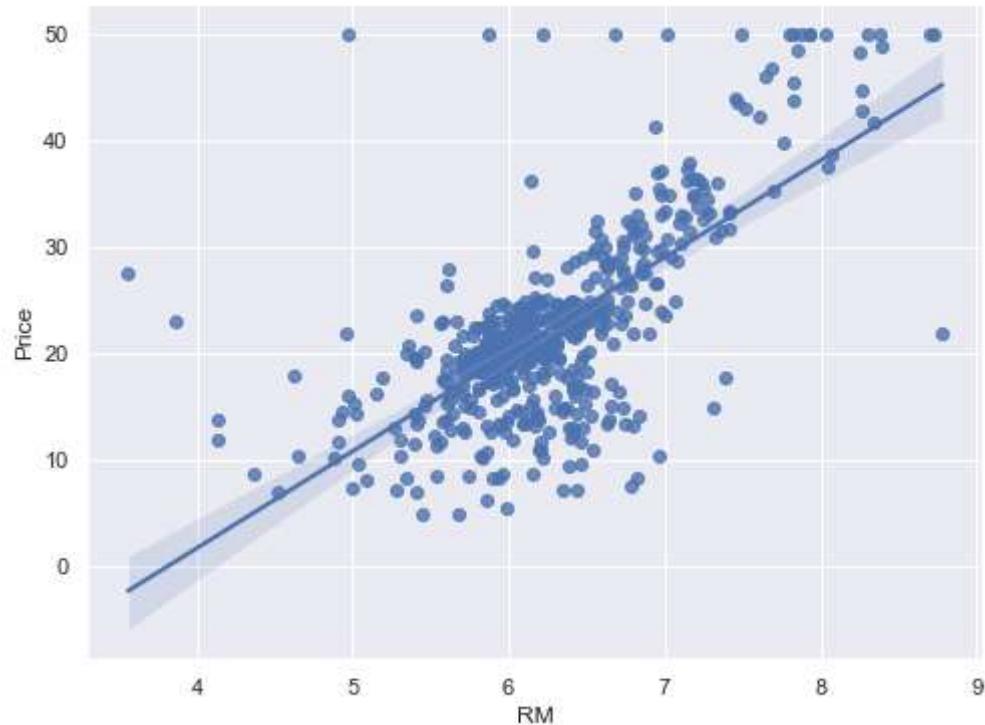
In [21]:

```
sns.set(rc={'figure.figsize':(8,6)})  
plt.scatter(df['RM'],df['Price']);
```



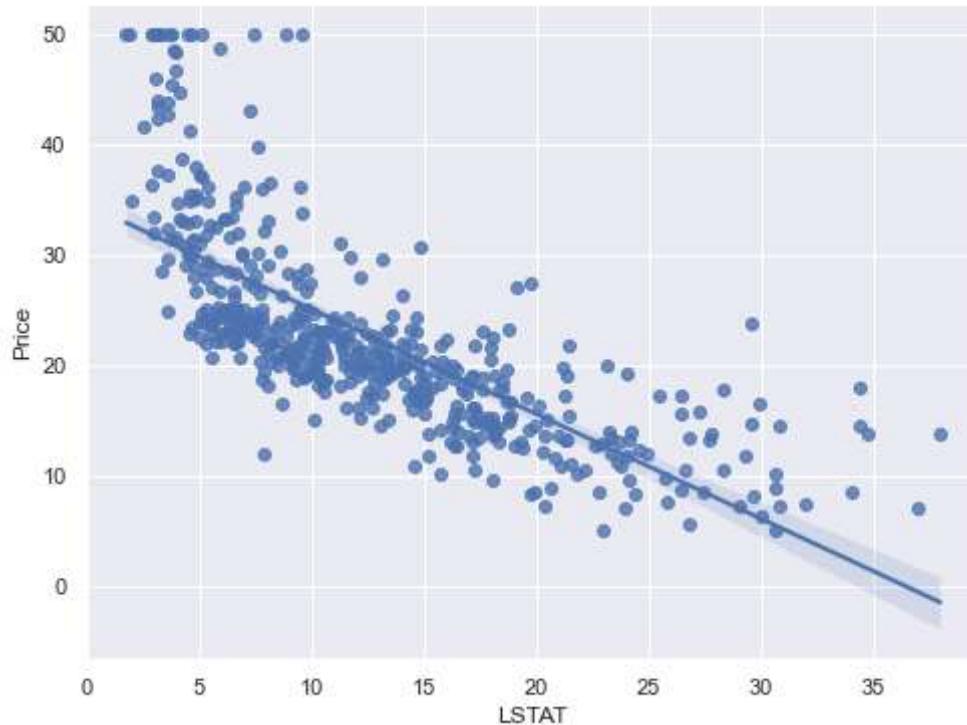
In [22]:

```
sns.regplot(x="RM",y="Price",data=df);
```



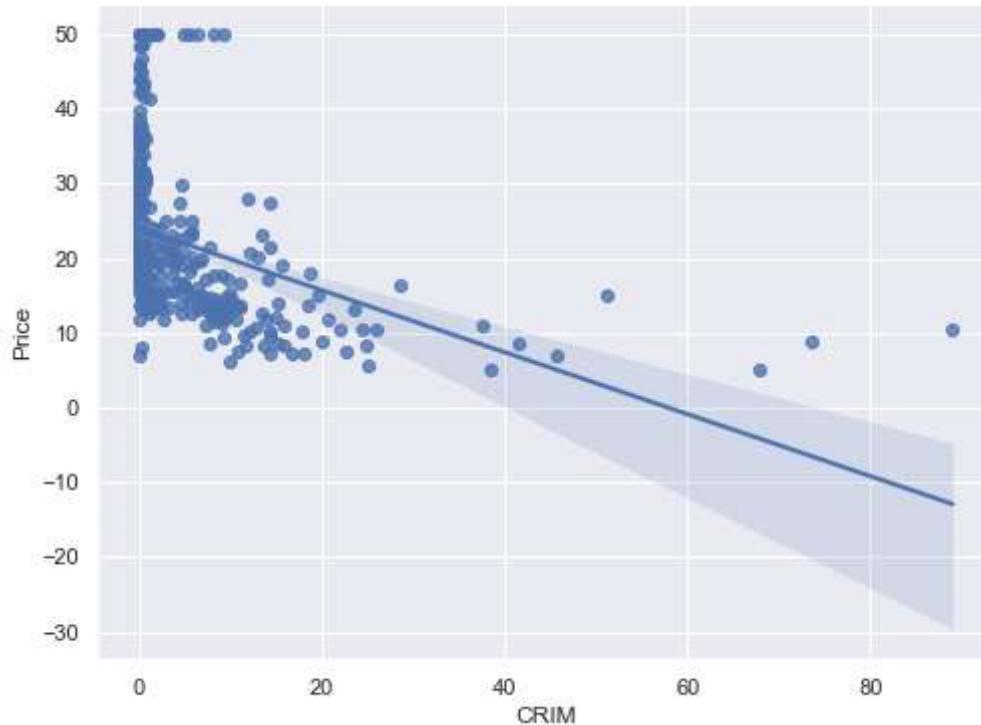
In [23]:

```
sns.regplot(x="LSTAT",y="Price",data=df);
```



In [24]:

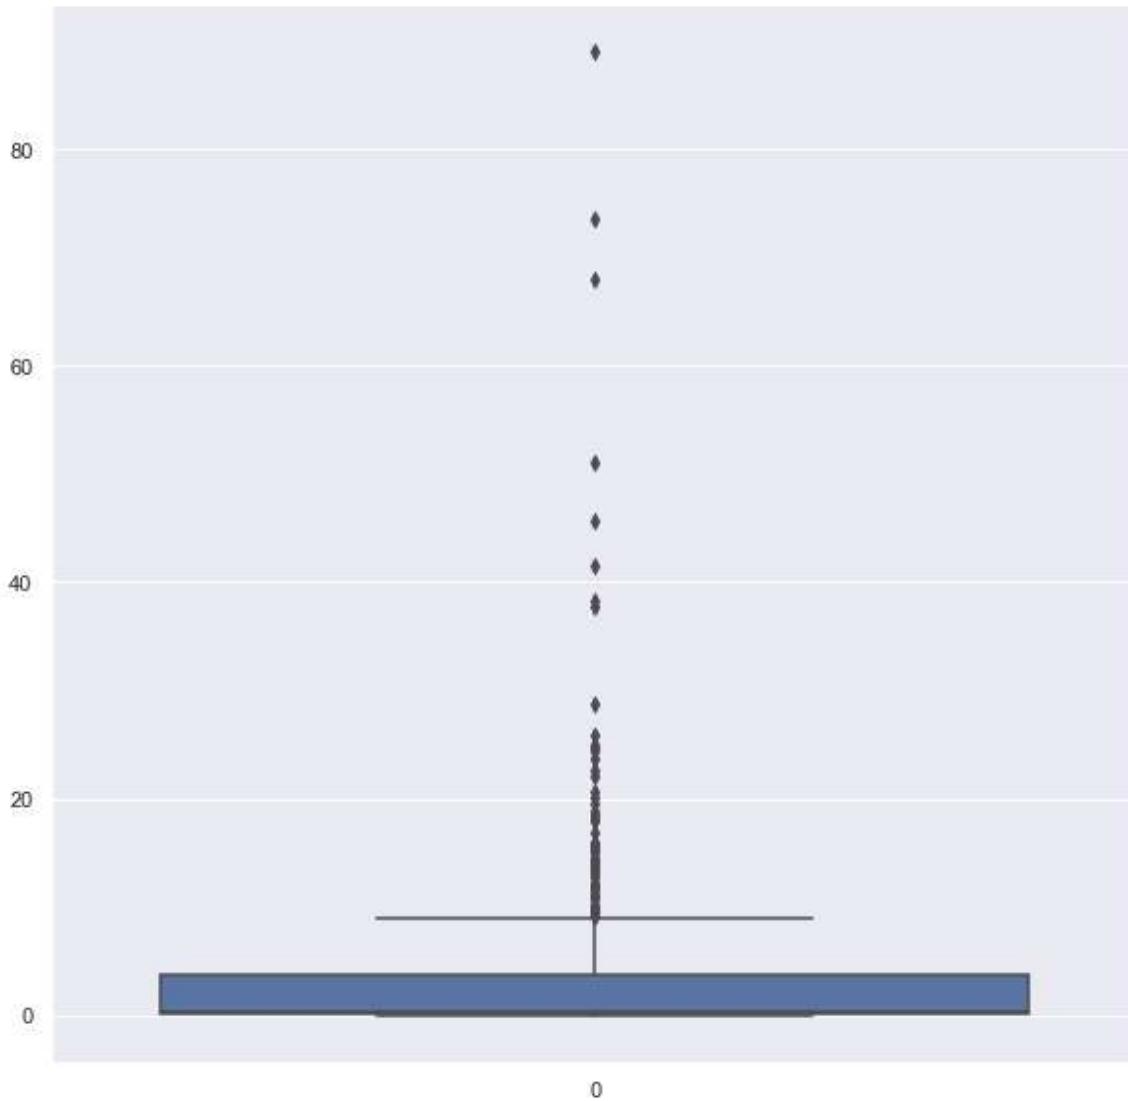
```
sns.regplot(x="CRIM",y="Price",data=df);
```



Detecting the Outliers present in the Criminal Rate & Price using Box-Plot

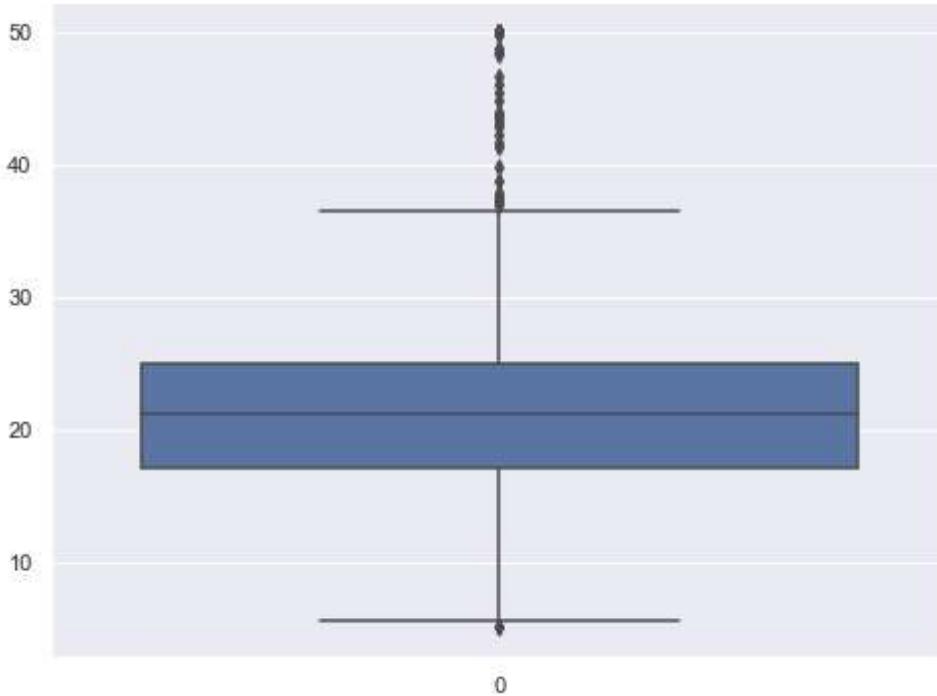
In [35]:

```
plt.figure(figsize=(10,10))
sns.boxplot(df['CRIM']);
```



In [36]:

```
sns.boxplot(df['Price']);
```



In [37]:

```
# Checking the DataFrame and first 5 rows
df.head()
```

Out[37]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90

Splitting the Dataset into Features and Label

In [38]:

```
# Independent and Dependent Features
```

```
X=df.iloc[:, :-1]
y=df.iloc[:, -1]
```

In [41]:

```
X
```

Out[41]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	E
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90
...
501	0.06263	0.0	11.93	0.0	0.573	6.593	69.1	2.4786	1.0	273.0	21.0	391.95
502	0.04527	0.0	11.93	0.0	0.573	6.120	76.7	2.2875	1.0	273.0	21.0	396.90
503	0.06076	0.0	11.93	0.0	0.573	6.976	91.0	2.1675	1.0	273.0	21.0	396.90
504	0.10959	0.0	11.93	0.0	0.573	6.794	89.3	2.3889	1.0	273.0	21.0	393.45
505	0.04741	0.0	11.93	0.0	0.573	6.030	80.8	2.5050	1.0	273.0	21.0	396.90

506 rows × 13 columns

In [42]:

```
y
```

Out[42]:

```
0      24.0
1      21.6
2      34.7
3      33.4
4      36.2
...
501    22.4
502    20.6
503    23.9
504    22.0
505    11.9
Name: Price, Length: 506, dtype: float64
```

Model Training

In [43]:

```
from sklearn.model_selection import train_test_split
```

Splitting the Dataset into Train & Test Data

In [44]:

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state=10)
```

In [45]:

```
X_train.shape
```

Out[45]:

```
(339, 13)
```

In [46]:

```
y_train.shape
```

Out[46]:

```
(339,)
```

In [47]:

```
X_test.shape
```

Out[47]:

```
(167, 13)
```

In [48]:

```
y_test.shape
```

Out[48]:

```
(167,)
```

Scaling the Train & Test Data Set using Standard Scaler

In [49]:

```
## Standardized the Dataset
from sklearn.preprocessing import StandardScaler
scaler=StandardScaler()
```

In [50]:

```
scaler
```

Out[50]:

```
StandardScaler()
```

In [51]:

```
X_train_scaled = scaler.fit_transform(X_train)
```

In [52]:

```
X_test_scaled = scaler.transform(X_test)
```

In [53]:

```
X_train_scaled
```

Out[53]:

```
array([[-0.13641471, -0.47928013,  1.16787606, ..., -1.77731527,
       0.39261401,  2.36597873],
      [-0.41777807, -0.47928013, -1.18043314, ..., -0.75987458,
       0.14721899, -0.54115799],
      [ 1.31269177, -0.47928013,  0.95517731, ...,  0.76628645,
       0.19334986,  2.52100705],
      ...,
      [-0.13520965, -0.47928013,  0.95517731, ...,  0.76628645,
       0.17012536,  0.06331026],
      [-0.40281114, -0.47928013,  2.04022838, ...,  0.25756611,
       0.32166792,  0.27238516],
      [-0.33104058,  0.34161649, -1.07552092, ..., -2.56351944,
       0.39993132, -0.34772815]])
```

In [54]:

```
X_test_scaled
```

Out[54]:

```
array([[-0.41664568,  0.87519929, -1.33277144, ..., -0.06616502,
       0.41011193, -0.56391444],
      [-0.42063267,  1.98340973, -1.22498491, ..., -1.36108953,
       0.41021798, -1.11860295],
      [-0.41894074,  2.80430634, -1.16175014, ..., -1.12985301,
       0.44765291, -1.16980497],
      ...,
      [-0.40804678,  1.36773726, -1.15169007, ..., -1.54607875,
       0.29854946, -1.18545003],
      [-0.41098494, -0.47928013,  0.19779729, ...,  0.07257689,
       0.20119741, -0.13154186],
      [-0.37856708, -0.47928013, -0.22328875, ..., -0.06616502,
       0.43482111, -0.5141347 ]])
```

Model Training

In [55]:

```
from sklearn.linear_model import LinearRegression
```

In [56]:

```
regression = LinearRegression()
```

In [57]:

```
regression
```

Out[57]:

```
LinearRegression()
```

In [58]:

```
regression.fit(X_train,y_train)
```

Out[58]:

```
LinearRegression()
```

In [59]:

```
## Print the coefficient and the intercept
print(regression.coef_)
```

```
[-1.56353750e-01  6.60616548e-02 -2.01655290e-02  1.57889806e+00
 -1.53490777e+01  3.34105533e+00  1.14773954e-02 -1.55360652e+00
  3.05450100e-01 -1.21583833e-02 -8.14405729e-01  1.32763069e-02
 -5.45290138e-01]
```

In [60]:

```
print(regression.intercept_)
```

```
34.21329725760603
```

Prediction for the Test data

In [61]:

```
reg_pred=regression.predict(X_test)
```

In [62]:

```
reg_pred
```

Out[62]:

```
array([31.43849583, 31.98794389, 30.99895559, 22.31396689, 18.89492791,
       16.21371128, 35.9881236 , 14.81264582, 25.04500847, 37.12806894,
       21.49110158, 30.88757187, 28.05752881, 34.05600093, 33.75791114,
       40.63880011, 24.24023412, 23.41351375, 25.54158122, 21.34135664,
       32.71699711, 17.88341061, 25.49549436, 25.01006418, 32.54102925,
       20.48979076, 19.48816948, 16.92733183, 38.38530857, 0.36265208,
       32.42715816, 32.15306983, 26.10323665, 23.79611814, 20.67497128,
       19.69393973, 3.50784614, 35.26259797, 27.04725425, 27.66164435,
       34.35132103, 29.83057837, 18.40939436, 31.56953795, 17.91877807,
       28.50042742, 19.49382421, 21.69553078, 38.0954563 , 16.44490081,
       24.58507284, 19.67889486, 24.53954813, 34.30610423, 26.74699088,
       34.87803562, 21.06219662, 19.87980936, 18.68725139, 24.71786624,
       19.96344041, 23.56002479, 39.57630226, 42.81994338, 30.37060855,
       17.03737245, 23.83719412, 3.2425022 , 31.5046382 , 28.63779884,
       18.49288659, 27.14115768, 19.67125483, 25.34222917, 25.05430467,
       10.29463949, 38.96369453, 8.26774249, 18.52214761, 30.34082002,
       22.87681099, 20.96680268, 20.04604103, 28.73415756, 30.81726786,
       28.23002473, 26.28588806, 31.59181918, 22.13093608, -6.48201197,
       21.53000756, 19.90826887, 24.96686716, 23.44746617, 19.28521216,
       18.75729874, 27.40013804, 22.17867402, 26.82972 , 23.39779064,
       23.9260607 , 19.16632572, 21.09732823, 11.01452286, 13.7692535 ,
       20.74596484, 23.54892211, 14.04445469, 28.88171403, 15.77611741,
       15.25195598, 22.429474 , 26.60737213, 28.88742175, 24.29797261,
       18.26839956, 16.26943281, 17.40100292, 15.53131616, 21.27868825,
       33.78464602, 30.00899396, 21.16115702, 13.95560661, 16.18475215,
       29.30998858, 13.1866784 , 22.08393725, 24.34499386, 31.86829501,
       33.45923602, 5.90671516, 35.20153265, 24.17614831, 17.54200544,
       24.25032915, 28.44671354, 34.50123773, 6.33164665, 1.93565618,
       28.40727267, 12.56461105, 18.31045646, 19.71015745, 5.50105857,
       14.51366874, 37.193992 , 25.81821367, 23.31632083, 26.43254504,
       11.38255141, 20.46224115, 35.27645709, 20.57841598, 11.48799917,
       16.23913171, 24.56511742, 10.53131603, 15.07115005, 25.98488217,
       11.2136222 , 11.695686 , 19.40437966, 19.58768384, 32.43800883,
       22.66170871, 25.68576052])
```

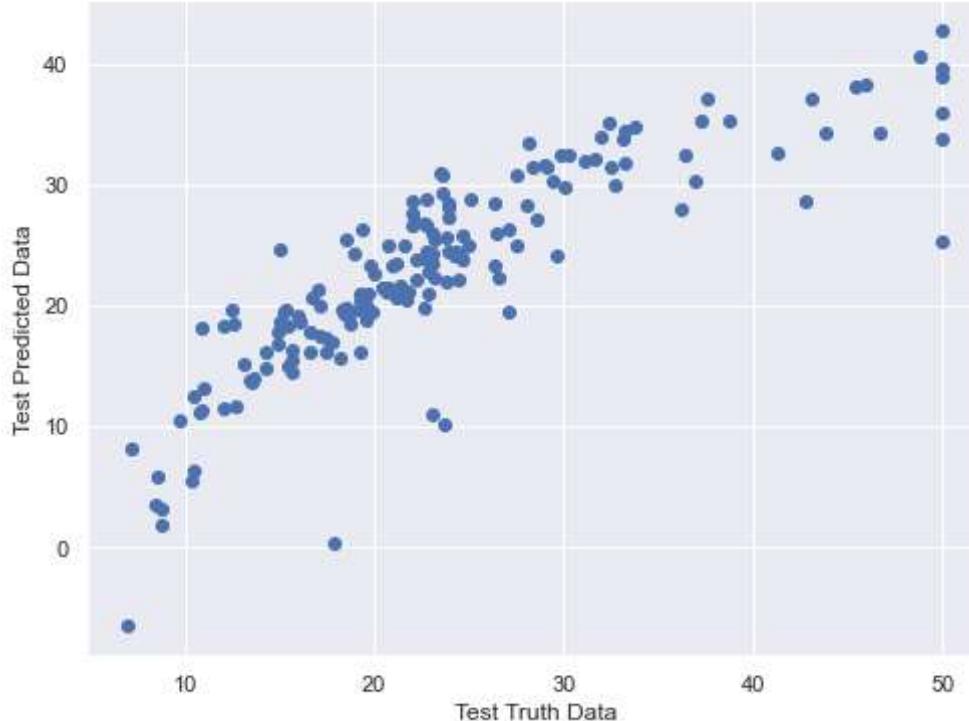
Assumptions of Linear Regression

In [63]:

```
plt.scatter(y_test,reg_pred)
plt.xlabel("Test Truth Data")
plt.ylabel("Test Predicted Data")
```

Out[63]:

```
Text(0, 0.5, 'Test Predicted Data')
```



Residuals

In [64]:

```
residuals = y_test-reg_pred
```

In [65]:

```
residuals
```

Out[65]:

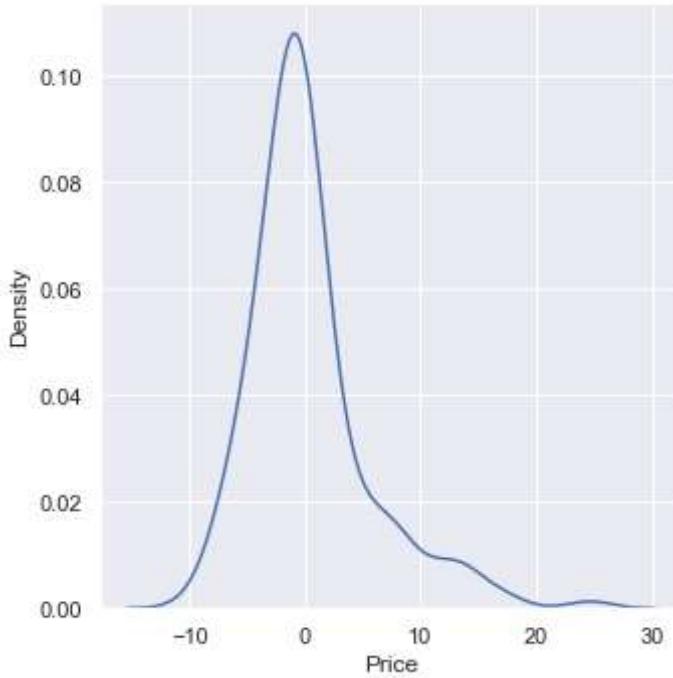
```
305    -3.038496
193    -0.887944
65     -7.498956
349     4.286033
151     0.705072
...
442    -1.004380
451    -4.387684
188    -2.638009
76     -2.661709
314    -1.885761
Name: Price, Length: 167, dtype: float64
```

In [66]:

```
sns.displot(residuals,kind='kde')
```

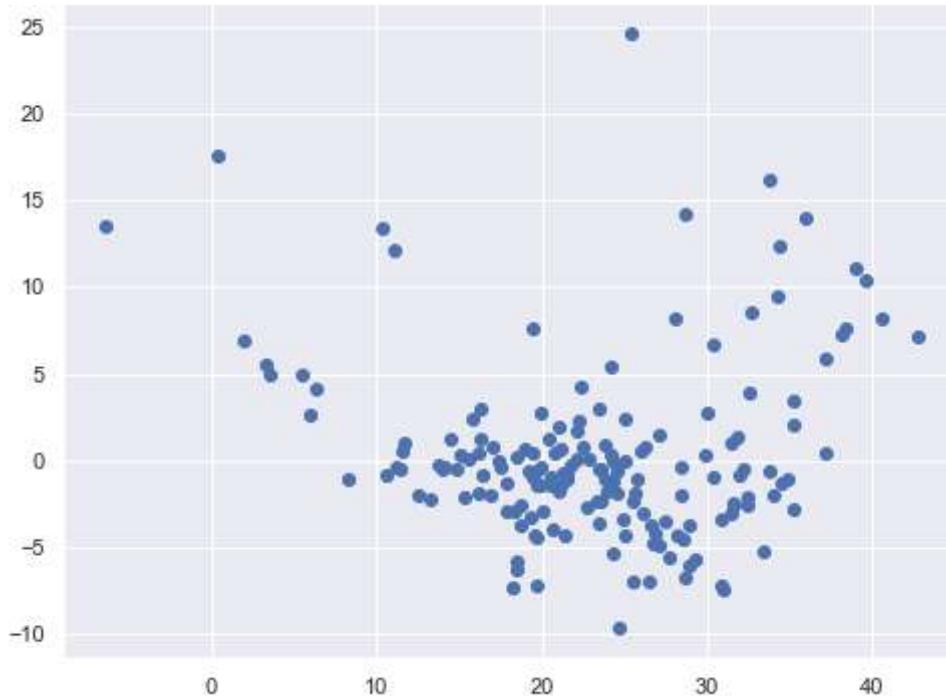
Out[66]:

```
<seaborn.axisgrid.FacetGrid at 0x1e742e3ee80>
```



In [67]:

```
## Scatter plot with predictions and residual  
## Uniform distribution  
plt.scatter(reg_pred,residuals);
```



Performance Metrics

In [68]:

```
from sklearn.metrics import mean_squared_error #MSE  
from sklearn.metrics import mean_absolute_error #MAE  
  
print(mean_squared_error(y_test,reg_pred))  
print(mean_absolute_error(y_test,reg_pred))  
print(np.sqrt(mean_squared_error(y_test,reg_pred)))
```

```
27.10099170996272  
3.5206585298798116  
5.205861284164487
```

R Square & Adjusted R square

In [69]:

```
from sklearn.metrics import r2_score
score = r2_score(y_test, reg_pred)
print(score)
```

0.716521939396753

In [70]:

```
## Adjusted R square
1 - (1-score)*(len(y_test)-1)/(len(y_test)-X_test.shape[1]-1)
```

Out[70]:

0.6924355682343857

Evaluating the Dataset Using Ridge Regression

In [71]:

```
from sklearn.linear_model import Ridge
ridge = Ridge()
```

In [72]:

ridge

Out[72]:

Ridge()

In [73]:

```
ridge.fit(X_train,y_train)
```

Out[73]:

Ridge()

In [74]:

```
print(ridge.coef_)
```

```
[-1.52632064e-01  6.72194813e-02 -4.61868347e-02  1.45844821e+00
 -7.76898482e+00  3.39133018e+00  4.68041987e-03 -1.43319434e+00
  2.87697853e-01 -1.27154968e-02 -7.34594385e-01  1.36842091e-02
 -5.54979507e-01]
```

In [75]:

```
print(ridge.intercept_)
```

28.878811583899946

In [76]:

```
ridge_pred = ridge.predict(X_test)
```

In [77]:

```
ridge_pred
```

Out[77]:

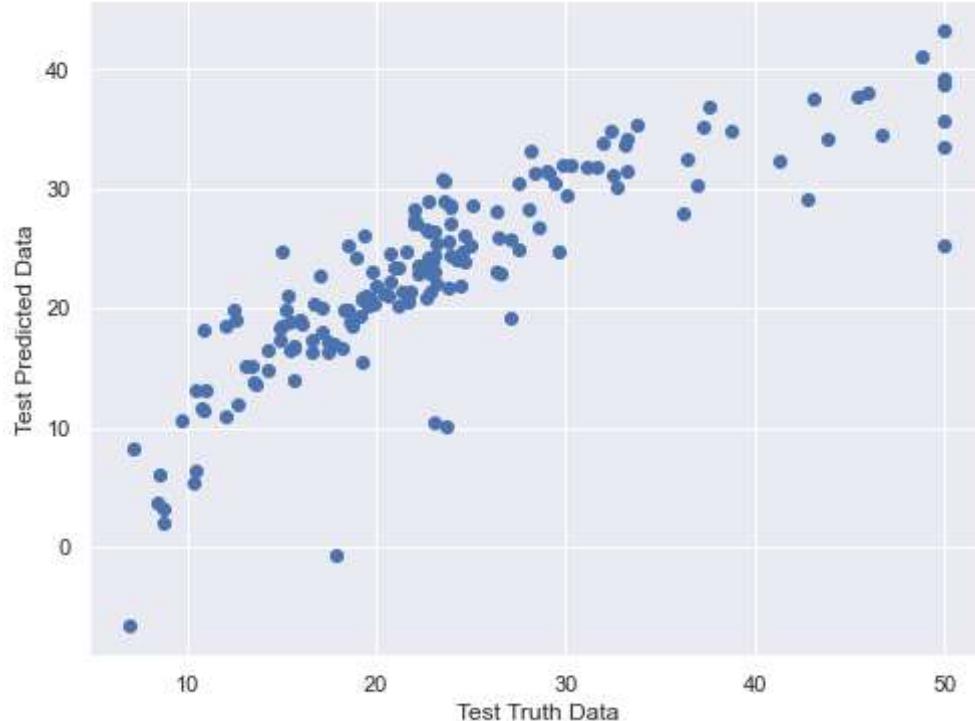
```
array([31.30151302, 31.90126709, 30.82061573, 22.88812126, 20.31302387,
       16.53738624, 35.72812342, 14.8816675 , 24.63531595, 36.860872 ,
       21.24748513, 30.46713945, 27.88750209, 33.88238471, 33.62448674,
       41.10851559, 24.75396381, 23.02945622, 25.51134988, 22.69511097,
       32.34032533, 18.39610936, 25.19196888, 25.25847581, 32.47842797,
       20.99995617, 19.2369491 , 17.40718641, 38.00318156, -0.61359473,
       32.04483047, 31.84921686, 26.43869186, 23.96961375, 20.30489338,
       19.93571226, 3.79628876, 35.24712422, 27.18060759, 27.45307204,
       34.54591726, 29.52160214, 18.50398038, 31.41198471, 17.43190724,
       28.64573998, 20.33141004, 21.39198549, 37.67651978, 16.60438885,
       24.20570719, 19.80900064, 24.26312305, 34.20633563, 27.07201815,
       35.29920299, 20.89256765, 20.83539778, 18.77362424, 24.77911964,
       20.14183389, 23.37554369, 39.23285572, 43.28676132, 30.45714887,
       17.07059688, 23.63831031, 3.2257982 , 31.19529737, 29.2191736 ,
       19.07024992, 26.74560227, 19.32823778, 25.20390851, 25.00399356,
       10.14495703, 38.69984714, 8.27960583, 18.53194663, 30.27261436,
       22.83546957, 21.45322051, 20.0093348 , 28.32764591, 30.62195844,
       28.53535251, 25.73580192, 31.53955632, 22.90414042, -6.6242741 ,
       22.28964733, 19.83165899, 24.71186029, 23.8500573 , 19.05308704,
       18.56224278, 27.08227103, 21.84161734, 26.54715488, 23.14882612,
       23.64448642, 18.91127471, 20.94359851, 10.47836918, 13.74582082,
       20.16243679, 23.05737754, 13.58013444, 28.93337707, 16.60270565,
       15.20197879, 22.03277272, 26.40083872, 28.57075083, 24.18673827,
       18.21939098, 15.54580923, 17.11763407, 16.8511474 , 21.08776724,
       33.44724064, 30.07125145, 21.46975178, 15.11884516, 16.40552026,
       28.95082498, 13.21512515, 21.71702584, 24.4563997 , 31.55015189,
       33.24347968, 6.08415598, 34.94161154, 24.69756072, 18.0754461 ,
       24.06272989, 28.18826271, 34.2155616 , 6.35685482, 2.00959625,
       28.36484928, 13.10018643, 18.82610551, 21.08244808, 5.40736266,
       14.02778007, 37.60352183, 26.13219721, 23.46037424, 26.08457696,
       11.54741208, 20.55527427, 34.91876361, 20.72969929, 10.96716828,
       16.38493095, 24.44222592, 10.61017709, 16.45530346, 26.02233123,
       11.61373636, 11.96109835, 19.91938337, 19.9460085 , 32.07539476,
       21.82189814, 25.68879093])
```

In [78]:

```
plt.scatter(y_test,ridge_pred)
plt.xlabel("Test Truth Data")
plt.ylabel("Test Predicted Data")
```

Out[78]:

Text(0, 0.5, 'Test Predicted Data')



Residuals

In [79]:

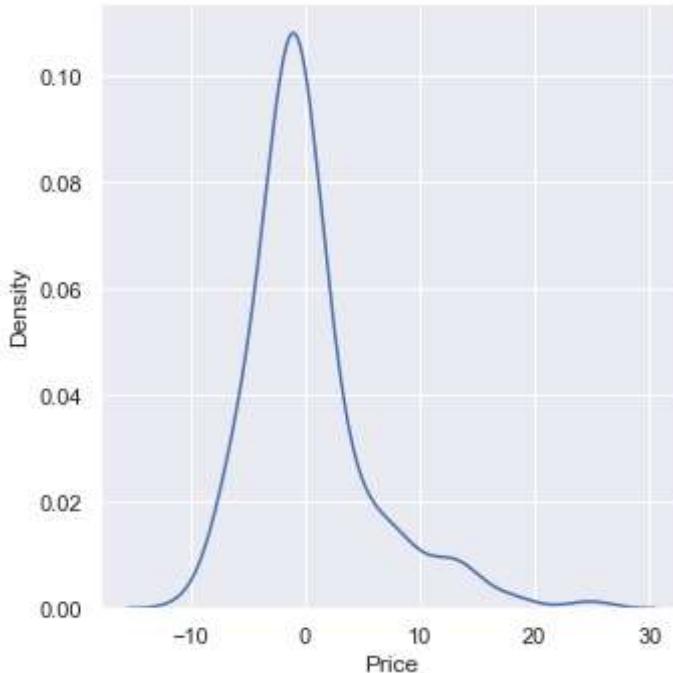
```
ridge_residuals = y_test - ridge_pred  
ridge_residuals
```

Out[79]:

```
305    -2.901513  
193    -0.801267  
65     -7.320616  
349     3.711879  
151    -0.713024  
      ...  
442    -1.519383  
451    -4.746009  
188    -2.275395  
76     -1.821898  
314    -1.888791  
Name: Price, Length: 167, dtype: float64
```

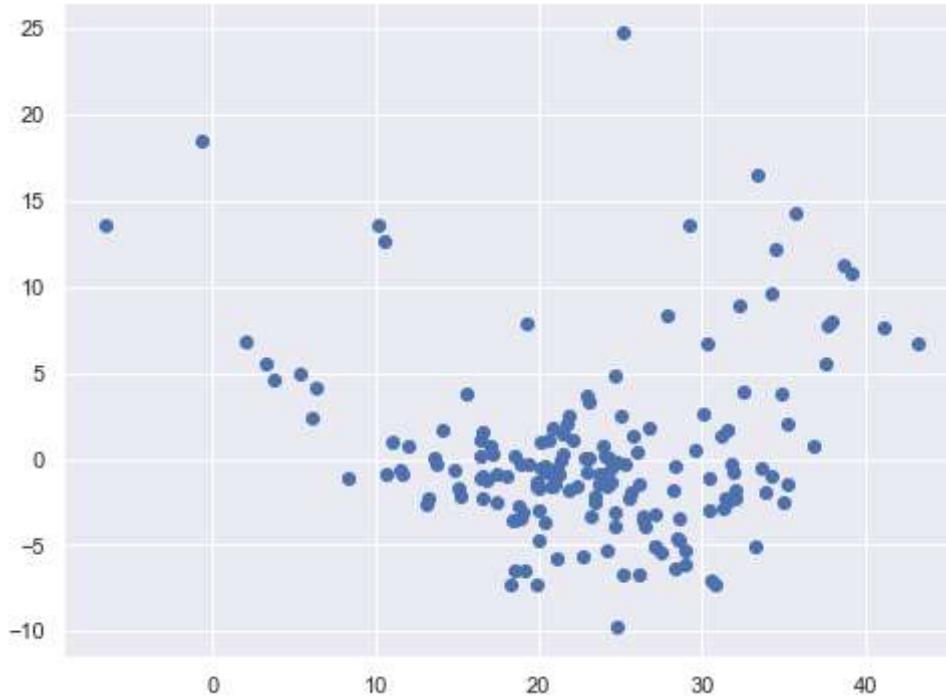
In [80]:

```
sns.displot(ridge_residuals, kind="kde");
```



In [81]:

```
plt.scatter(ridge_pred,ridge_residuals);
```



In [82]:

```
from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
print(mean_squared_error(y_test,ridge_pred))
print(mean_absolute_error(y_test,ridge_pred))
print(np.sqrt(mean_squared_error(y_test,ridge_pred)))
```

27.63979628041237

3.566042008593924

5.257356396556388

R square & Adjusted Rsquare

In [83]:

```
from sklearn.metrics import r2_score
ridge_score=r2_score(y_test,ridge_pred)
print(ridge_score)
```

0.7108860100436925

In [84]:

```
## Adjusted R square  
1 - (1-ridge_score)*(len(y_test)-1)/(len(y_test)-X_test.shape[1]-1)
```

Out[84]:

0.6863207690670128

Evaluating the dataset using Lasso Regression

In [85]:

```
from sklearn.linear_model import Lasso  
lasso = Lasso()
```

In [86]:

lasso

Out[86]:

Lasso()

In [87]:

```
lasso.fit(X_train,y_train)
```

Out[87]:

Lasso()

In [88]:

```
print(lasso.intercept_)
```

41.983204728562384

In [89]:

```
print(lasso.coef_)
```

```
[-0.10475563  0.07649297 -0.01718277  0.          -0.          0.24993515  
 0.03023573 -0.80122944  0.26755925 -0.01537057 -0.58410577  0.01011944  
-0.76343344]
```

In [90]:

```
lasso_pred=lasso.predict(X_test)
```

In [91]:

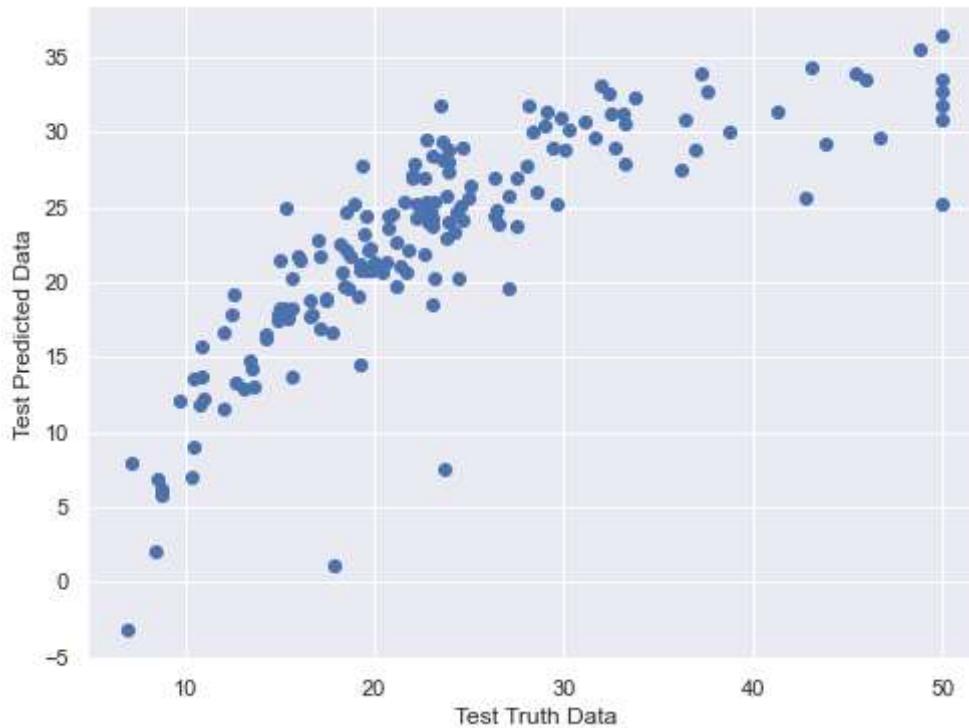
lasso_pred

Out[91]:

```
array([30.05062227, 30.75863386, 31.79136406, 23.81132585, 24.46450905,
       16.51502359, 33.55687901, 16.18590265, 23.63422119, 32.67378583,
       20.62340113, 26.88616937, 27.50950406, 33.13800373, 31.26216912,
       35.51177715, 25.1322652 , 24.37702007, 25.3426104 , 22.79399868,
       31.42535013, 17.87518719, 24.62140298, 25.67592917, 30.8570178 ,
       23.19651817, 19.54532952, 17.48554848, 33.51934661, 1.0746952 ,
       30.2290915 , 29.6311516 , 28.3746209 , 24.09051453, 17.8996287 ,
       20.68626471, 2.03311885, 33.92194066, 27.84246939, 26.98859872,
       29.66201241, 28.76027267, 16.66430744, 31.3070688 , 17.76180501,
       28.8142132 , 20.84880796, 21.08094697, 33.95869027, 18.19713582,
       25.3790762 , 17.83141295, 24.62557761, 29.23048683, 27.09040407,
       32.30973738, 21.12865331, 21.8682558 , 21.45083035, 21.4927721 ,
       20.82460129, 22.70209723, 32.64112207, 36.46204926, 29.01296777,
       16.60412861, 24.21223336, 6.15819233, 31.29938975, 25.66665157,
       19.12096902, 26.07635104, 19.09894122, 25.20532246, 23.72293349,
       7.46519487, 31.78184104, 7.93417039, 21.66650902, 28.82264119,
       24.04887629, 23.99760225, 21.73647125, 27.1081459 , 29.34325776,
       28.04599659, 25.75769079, 30.4155572 , 25.14780952, -3.2483763 ,
       24.33980236, 22.19039081, 25.32648911, 24.73588935, 21.71166789,
       18.17949212, 27.40772395, 20.30800437, 26.99320873, 22.2873327 ,
       24.23833743, 19.64886888, 22.15222613, 18.57341744, 14.18514138,
       19.70755856, 23.78098739, 13.03503894, 29.45659639, 22.47550945,
       12.90394483, 20.25757185, 24.61575068, 26.37500077, 25.18758983,
       15.73313472, 14.49559824, 18.87614435, 20.19492705, 21.29726476,
       30.81860198, 28.94638256, 22.17757046, 14.71215667, 18.79736922,
       28.12286666, 12.19429838, 22.97031693, 24.30773588, 27.93884477,
       31.73300601, 6.81523194, 32.55092938, 25.22113859, 16.92807435,
       23.27550401, 26.92401426, 30.59006878, 8.93970768, 5.84737618,
       27.80566228, 13.58229434, 17.59584967, 24.93011863, 6.99103487,
       13.70103998, 34.3311626 , 29.01410928, 24.59277709, 27.81462485,
       13.65501713, 20.71519689, 30.01636162, 20.82701755, 11.58795903,
       18.72787631, 23.99187497, 12.02982796, 18.1481676 , 24.76157021,
       11.85369437, 13.24565577, 19.77471126, 18.29296243, 30.94733281,
       21.26773868, 25.73858979])
```

In [92]:

```
plt.scatter(y_test,lasso_pred)
plt.xlabel("Test Truth Data")
plt.ylabel("Test Predicted Data");
```



Residuals

In [93]:

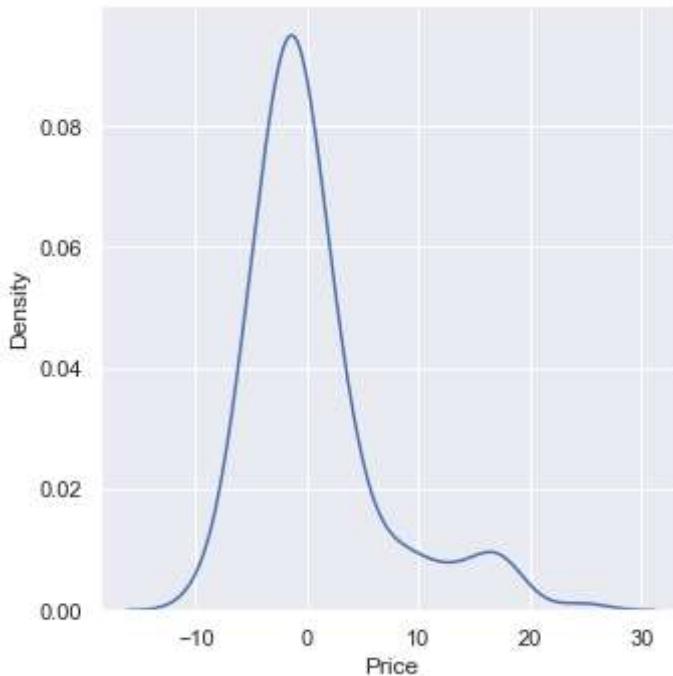
```
lasso_residuals = y_test-lasso_pred  
lasso_residuals
```

Out[93]:

```
305    -1.650622  
193     0.341366  
65     -8.291364  
349     2.788674  
151    -4.864509  
      ...  
442    -1.374711  
451    -3.092962  
188    -1.147333  
76     -1.267739  
314    -1.938590  
Name: Price, Length: 167, dtype: float64
```

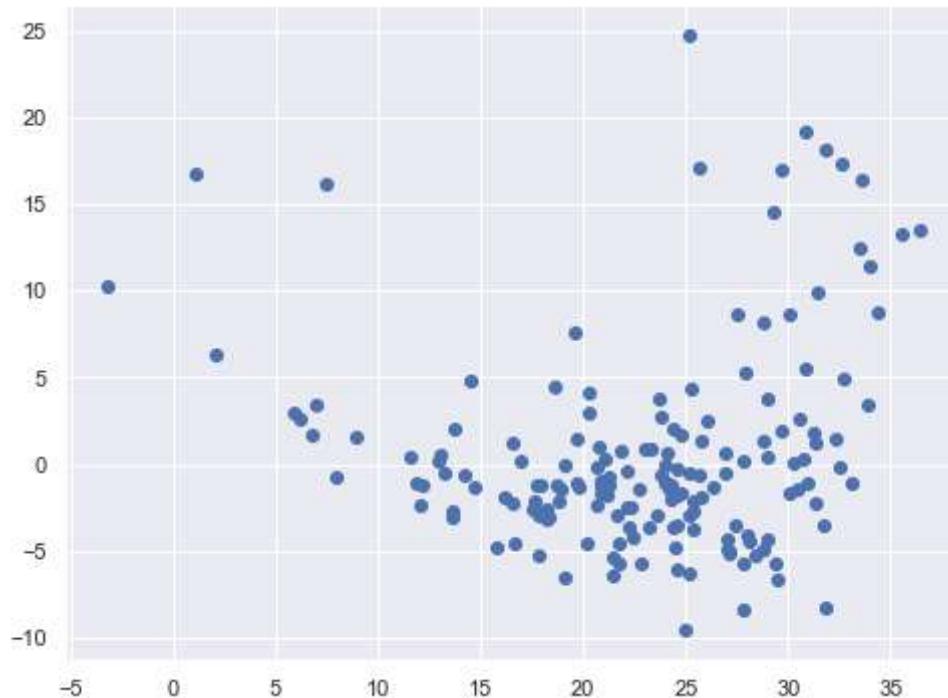
In [94]:

```
sns.displot(lasso_residuals,kind="kde");
```



In [96]:

```
## Scatter plot with predictions and residual
##Uniform distribution
plt.scatter(lasso_pred,lasso_residuals);
```



In [97]:

```
from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
print(mean_squared_error(y_test,lasso_pred))
print(mean_absolute_error(y_test,lasso_pred))
print(np.sqrt(mean_squared_error(y_test,lasso_pred)))
```

```
36.11332777156728
4.067062767559684
6.009436560241507
```

R Square & Adjusted R squared

In [98]:

```
from sklearn.metrics import r2_score
lasso_score=r2_score(y_test,lasso_pred)
print(lasso_score)
```

```
0.6222523430812359
```

In [99]:

```
## Adjusted R square  
#display adjusted R-squared  
1 - (1-lasso_score)*(len(y_test)-1)/(len(y_test)-X_test.shape[1]-1)
```

Out[99]:

```
0.5901561369378114
```

Evaluating the dataset with Elastic Net Regression

In [100]:

```
from sklearn.linear_model import ElasticNet  
elasticnet = ElasticNet()
```

In [101]:

```
elasticnet
```

Out[101]:

```
ElasticNet()
```

In [102]:

```
elasticnet.fit(X_train,y_train)
```

Out[102]:

```
ElasticNet()
```

In [103]:

```
print(elasticnet.intercept_)
```

```
39.51456077615768
```

In [104]:

```
print(elasticnet.coef_)
```

```
[-0.12157848  0.07459975 -0.03899046  0.          -0.          0.70123205  
 0.02893353 -0.84498184  0.28961626 -0.01546206 -0.61704509  0.01082719  
-0.73127687]
```

In [105]:

```
elasticnet_pred=elasticnet.predict(X_test)
```

In [106]:

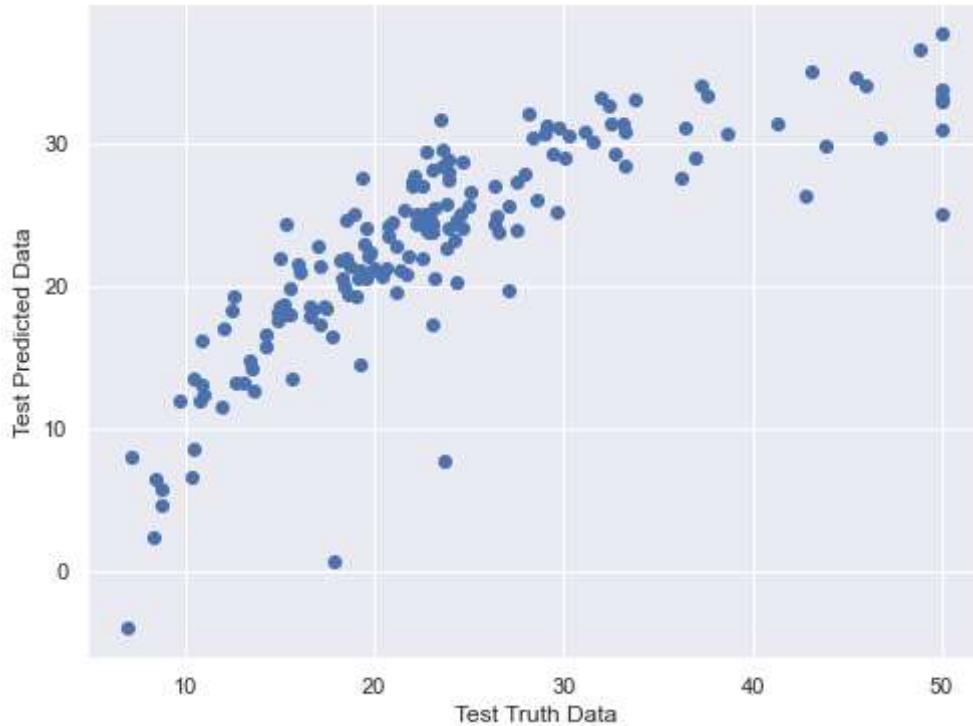
```
elasticnet_pred
```

Out[106]:

```
array([30.32727993, 30.82005104, 31.58152723, 23.74897548, 24.04066339,
       16.64208335, 33.72175392, 15.7796708 , 23.54360404, 33.38864033,
       20.64135992, 27.30279352, 27.61114038, 33.16496835, 31.30632428,
       36.60242063, 25.02604143, 24.32320649, 25.42409477, 22.74163791,
       31.42753395, 18.15172597, 24.66377356, 25.5930192 , 31.0902106 ,
       22.99406956, 19.75673457, 17.62225051, 34.05998978, 0.71393175,
       30.54158704, 30.05134167, 28.13413609, 23.98640017, 18.3419393 ,
       20.59655372, 2.36980673, 33.98725736, 27.65073017, 27.07927693,
       30.32046718, 28.98490321, 17.05004568, 31.23586807, 17.83741534,
       28.86472005, 21.03042701, 21.08354432, 34.61346048, 18.06846239,
       25.10524461, 18.26828696, 24.68020774, 29.84488335, 26.99724854,
       33.00363459, 21.04127003, 21.96291557, 20.9156221 , 21.93093756,
       20.51191466, 22.75831806, 33.25356336, 37.69317773, 29.21343037,
       16.44883562, 23.88399261, 5.73678016, 31.41397208, 26.2387969 ,
       19.29642893, 26.03173494, 19.20268502, 25.02297417, 23.84443518,
       7.68082768, 32.86644065, 8.05404899, 21.42877667, 29.02221506,
       24.02748648, 23.72705434, 21.43227946, 27.30244657, 29.59030615,
       28.01870181, 25.66336436, 30.59211532, 24.96621901, -3.98999389,
       24.18939757, 21.92147658, 25.29138706, 24.59227886, 21.58691263,
       18.50577433, 27.393029 , 20.27186886, 26.99600445, 22.37968176,
       24.27900606, 19.42893094, 22.0778287 , 17.35434124, 14.17049383,
       19.60954073, 23.7971747 , 12.6676568 , 29.38676612, 21.73985069,
       13.18186575, 20.55801528, 24.84735541, 26.63076152, 25.07411166,
       16.12276859, 14.45382955, 18.57745942, 19.85077271, 21.25627948,
       31.00768345, 29.21525794, 22.02696279, 14.74351416, 18.55224908,
       28.34160315, 12.37902394, 22.62806951, 24.32172973, 28.42443954,
       32.06470207, 6.48450097, 32.57411341, 25.20463016, 17.27912488,
       23.14124551, 27.06784434, 30.86575934, 8.5643464 , 4.70402049,
       27.82665403, 13.5627324 , 17.9576227 , 24.31428469, 6.64669939,
       13.52691815, 35.07905289, 28.71229522, 24.42196943, 27.53944472,
       13.05632256, 20.78080524, 30.71452488, 20.60297639, 11.49834979,
       18.48955421, 24.06063117, 11.95041588, 18.07009953, 24.8600991 ,
       11.94641703, 13.19510438, 20.02129294, 18.73517709, 31.02814718,
       21.18425857, 25.77300816])
```

In [108]:

```
plt.scatter(y_test,elasticnet_pred)
plt.xlabel("Test Truth Data")
plt.ylabel("Test Predicted Data");
```



Residuals

In [109]:

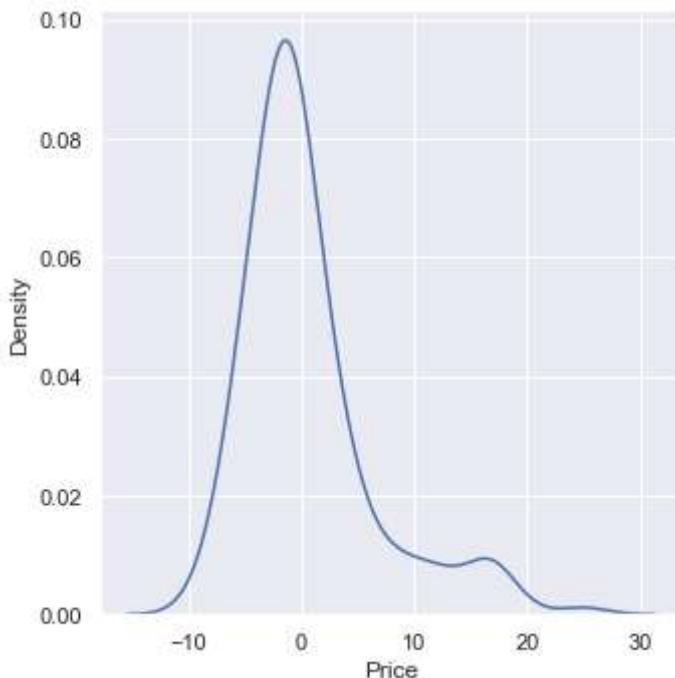
```
elasticnet_residuals=y_test-elasticnet_pred  
elasticnet_residuals
```

Out[109]:

```
305    -1.927280  
193     0.279949  
65     -8.081527  
349     2.851025  
151    -4.440663  
      ...  
442    -1.621293  
451    -3.535177  
188    -1.228147  
76     -1.184259  
314    -1.973008  
Name: Price, Length: 167, dtype: float64
```

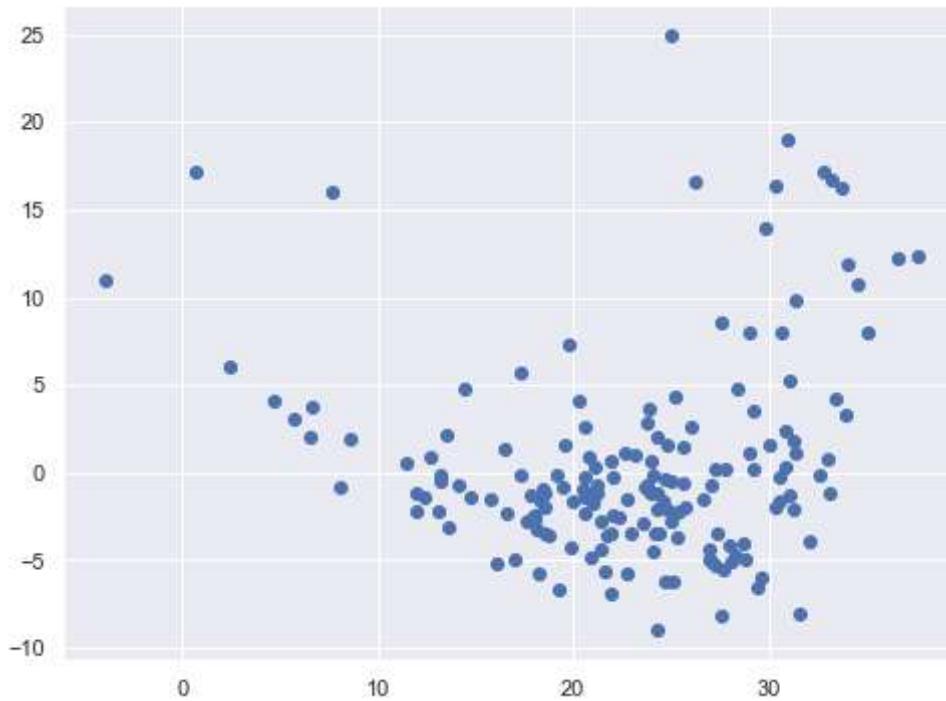
In [110]:

```
sns.displot(elasticnet_residuals,kind="kde");
```



In [111]:

```
## Scatter plot with predictions and residual  
##Uniform distribution  
plt.scatter(elasticnet_pred,elasticnet_residuals);
```



In [112]:

```
from sklearn.metrics import mean_squared_error  
from sklearn.metrics import mean_absolute_error  
print(mean_squared_error(y_test,elasticnet_pred))  
print(mean_absolute_error(y_test,elasticnet_pred))  
print(np.sqrt(mean_squared_error(y_test,elasticnet_pred)))
```

```
34.65454968329032  
4.002630706596711  
5.886811503971426
```

Rsquare & Adjusted R square

In [113]:

```
from sklearn.metrics import r2_score
elasticnet_score=r2_score(y_test,elasticnet_pred)
print(elasticnet_score)
```

0.6375112527086368

In [114]:

```
## Adjusted R square
#display adjusted R-squared
1 - (1-elasticnet_score)*(len(y_test)-1)/(len(y_test)-X_test.shape[1]-1)
```

Out[114]:

0.6067115552263642

Implemented Linear, Lasso, Ridge & Elastic-Net Regresion on Boston Housing Dataset.